

This work is licensed under a Creative Commons license



Attribution-NonCommercial-NoDerivatives 4.0 International  
(CC BY-NC-ND 4.0)

You are free to:

**Share** copy and redistribute the material in any medium or format.

Under the following terms:

**Attribution** You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.

**NonCommercial** You may not use the material for commercial purposes.

**NoDerivatives** If you remix, transform, or build upon the material, you may not distribute the modified material.

# Software security

Fantastic vulnerabilities and where to find them

Giovanni Lagorio

`giovanni.lagorio@unige.it`

`https://csec.it/people/giovanni\_lagorio`

Twitter & GitHub: `zxgio`

DIBRIS - Dipartimento di Informatica, Bioingegneria, Robotica e Ingegneria dei Sistemi  
University of Genova, Italy



`www.zenhack.it`

## 1 Introduction

## 2 Common Weaknesses

- Integer overflows (& undefined/implementation-defined behaviours)
- Command injection
- Memory errors
- “Hidden inputs”
- Race conditions
- Privileges
- Information leakage

## 3 Secure Coding

- Static Analysis
- Dynamic Analysis
  - Sanitizers
  - Fuzzers

# Terminology

**Bug/ flaw** A coding/design error; basically, defects

**Vulnerability** A **defect** that can be **exploited**, resulting in a negative **impact** to some **security** properties (e.g., confidentiality)

**Exploitation** The act of causing an **unintended behavior**, by taking advantage of vulnerabilities

**Exploit** A piece of software (that exploits “something”); typically categorized on how they connect to vulnerable software:

- a **remote exploit** works over a network and works without any prior access to the vulnerable system
- a **local exploit** requires prior access to the vulnerable system, and usually increases the privileges

# Vulnerability classification: CVE vs CWE

- **Specific instances**, within a product or system, of known vulnerabilities **are collected as CVE**, Common Vulnerabilities and Exposures — <https://cve.org/>
  - E.g., *CVE-2020-9641*  
Adobe Illustrator v24.1.2 and earlier have a memory corruption vulnerability. Exploitation could lead to arbitrary code execution.
  - This site does *not* provide proof of concepts/exploits; a repository for exploits and PoC, rather than advisories, is <https://www.exploit-db.com>
- Common **classes of vulnerabilities**, not specific to products/systems, **are collected as CWE**, Common Weakness Enumeration — <https://cwe.mitre.org/>
  - E.g., *CWE-242: Use of Inherently Dangerous Function*  
Program calls a function that can never be guaranteed to work safely.

A good, yet dated, book is *24 Deadly Sins of Software Security* [HLV10]

*Software with security vulnerabilities can be written in any programming language. Still, the programming language can make a difference here, by the language features it provides (or omits), and by the programmer support it offers for these, in the form of compilers, type checkers, run-time execution engines [Pol17]*

For instance, C and C++ have a lot of *undefined* behaviours [WCC<sup>+</sup>12]  
E.g., what does it happen when you divide an `int` by 0?

A very interesting talk is “Garbage In, Garbage Out: Arguing about Undefined Behavior” by Chandler Carruth @ CppCon 2016

[https://youtu.be/yG10Z69H\\_o](https://youtu.be/yG10Z69H_o)

For more, see [DMS06]

# Programming languages (updated: November 2022)

Which are the most popular/used today? (not necessarily the best)

- Python
- C/C++
- Java/Visual Basic.NET/C#
- PHP/Javascript/R
- Assembly in 8<sup>th</sup> position (!)

See, for instance:

- <https://www.tiobe.com/tiobe-index/> ... popularity ... based on the number of skilled engineers world-wide, courses and third party vendors. Popular search engines ... *is not about the best* ...
- <https://spectrum.ieee.org/top-programming-languages-2022>

# A common misconception

For a program to be secure, all its portions must be secure, not just the parts that explicitly address security (e.g., access control mechanisms)

- security features are typically implemented with the idea that they must function correctly to maintain system security
- like a chain, a program is only as strong as its weakest link
- non-security features are often the ones that can go wrong and lead to security problems



# Correctness $\neq$ Security

(non-security) Testing is not enough:

- Testing typically means running down the list of requirements, making sure they are fulfilled
- If the software fails to meet a particular requirement  $\rightarrow$  bug
- Security problems are frequently “unintended functionalities”, not requirement violations

quoting Ivan Arce, CTO of Core Security Technologies

*Reliable software does what it is supposed to do. Secure software does what it is supposed to do, and nothing else.*

# Ingredients for secure coding

- Write defect-free code 😊
- Perform **regular code reviews**
  - This **requires knowledge**; practitioners with little experience: do not know what to look for, or what to do about problems
    - A very good reference on security assessment is [DMS06]
    - An effective way to improve is **studying past errors** to prevent them from happening again
- Leverage **static analysis** tools [CW07] to
  - quickly check many possibilities and corner cases
  - explore a large number of “what if” scenarios, *without* running your (incomplete?) code
- Discover remaining vulnerabilities **dynamically, by leveraging**:
  - **fuzzers**
  - **sanitizers**
  - ...

# Secure coding != Defensive programming

## Defensive programming

Coding with the mindset that errors are inevitable and that, sooner or later, something will go wrong and lead to unexpected conditions.

Kernighan and Plauger [KP81] call it “writing the program so it can cope with small disasters”

- Expecting anomalies is a step in the right direction
  - However, defensive programming alone does *not* guarantee security
- We must assume the existence of an **adversary**, someone who is **intentionally trying to subvert** the system
  - Software security is about creating programs that behave correctly even in the presence of malicious behavior

For instance, ...

# Defensive programming example

Consider the following snippet:

```
void printMsg(FILE* file, char* msg)
{
    fprintf(file, msg);
}
```

What does it happen if file, or msg, is 0? With defensive programming:

```
void printMsg(FILE* file, char* msg)
{
    if (file == NULL)
        logErrorAndExit("attempt to print message to null file");
    else if (msg == NULL)
        logErrorAndExit("attempt to print null message");
    fprintf(file, msg);
}
```

Yet, this is still extremely insecure!

# Computer integers

Tricky point: what we typically call “integers” are not  $\mathbb{Z}$ , but  $\mathbb{Z}_{2^n}$

## Quick quiz

Suppose  $n = 3$ , so  $2^n = 8$

- In  $\mathbb{Z}_8$ ,  $2+3=$
- In  $\mathbb{Z}_8$ ,  $6+3=$
- Can you find  $x, y \in \mathbb{Z}_8 \setminus \{0\}$  s.t.  $x \cdot y = 0$ ?

Moreover,  $-1 = 7$ ,  $-2 = 6$ , and so on.

Typically,  $n$  bits are used to represent either the interval

- $[0, 2^n - 1]$  or
- $[-2^{n-1}, 2^{n-1} - 1]$

# Truncation and extension

In C/C++, what happens when the assignment `a=b` is executed? Or when the expression `a+b` is evaluated?

It depends on:

- `sizeof(a)` and `sizeof(b)`
- the “signedness” of `a` and `b`

... and can be rather tricky: generally, mixing signed and unsigned yield unsigned.

When `sizeof(a) != sizeof(b)` there is truncation or 0/sign-extension

In [DMS06] you can find about 100 pages on C issues!

# Truncation/extension example

```
#include <stdio.h>

int main()
{
    int i = 0xcafebabe; /* in this example sizeof(int) is 4 */
    short s = i;
    unsigned short us = i;
    signed char c = i;
    unsigned char uc = i;
    printf("i=%x\n", i);
    printf("s=%x us=%x\n", (int)s, (int)us);
    printf("c=%x uc=%x\n", (int)c, (int)uc);
}
```

the result is:

```
i=cafebabe
s=ffffbabe us=babe
c=ffffffbe uc=be
```

# From the C++11 standard [ANS12]

## Section 3.9.1

Unsigned integers ... shall obey the laws of arithmetic modulo  $2^n$  where  $n$  is the number of bits in the value representation of that particular size of integer.

A footnote clarifies that: “This implies that unsigned arithmetic does not overflow ...”

## Section 5 [expr]

If ... the result is ... not in the range of representable values for its type, the **behavior is undefined**



# Bound checking (?)

```
void f(int i, int j) {  
    int k;  
    if (i<0 || j<0) {  
        printf("i and j must be non-negative!\n");  
        return; }  
    k = i + j;  
    if (k >= 100) {  
        printf("i and j are too big!\n");  
        return; }  
    printf("%d is in the interval [0, 99]\n", k);  
    /* (***) use k; e.g. to index an array of 100 elements */  
}
```

## Quick quiz

Can you find  $i$  and  $j$  s.t.  $f(i, j)$  reaches (\*\*\*) but  $k \notin [0, 99]$ ?

# Arithmetic checking (?)

```
int f(int a, int b, int c) {  
    if (b==0 || c==0) {  
        printf("b and c cannot be zero!\n");  
        return 0;  
    }  
    return a / (b*c);  
}
```

## Quick quiz

Can you find three arguments that make `f` crash?

For instance, on *some* systems

Floating point exception (core dumped)

# Allocation checking (?)

```
int *array_copy(int *array, int len) {  
    if (array==0 || len<0)  
        return 0;  
    int *copy, i;  
    copy = malloc(len * sizeof(int));  
    if (copy == 0)  
        return 0;  
    for(i = 0; i < len; i++)  
        copy[i] = array[i];  
    return copy  
}
```

## Quick quiz

What can go wrong here?

# Real-world example 1: OpenSSH 2.3.1–3.3 (2002)

*Several versions of OpenSSH's sshd between 2.3.1 and 3.3 contain an input validation error that can result in an **integer overflow and privilege escalation**.*

- <https://www.kb.cert.org/vuls/id/369347>
- <https://www.openssh.com/txt/preauth.adv>

## Real-world example 2: Stagefright (July, 2015)



Stagefright is the collective name for a group of software bugs that affect Android operating system, **allowing an attacker to perform arbitrary operations on the victim device** through remote code execution and privilege escalation.

Security researchers demonstrate the bugs with a proof of concept that sends specially crafted MMS messages to the victim device and in most cases **requires no end-user actions upon message reception** to succeed, while using the phone number as the only target information.

The underlying attack vector exploits certain **integer overflow vulnerabilities**

<https://www.kb.cert.org/vuls/id/924951>

## Real world example 3: Boeing 787 (2020)

The US Federal Aviation Administration has ordered Boeing 787 operators to **switch their aircraft off and on every 51 days to prevent what it called “several potentially catastrophic failure scenarios”** — including the crashing of onboard network switches. ...

According to the directive itself, if the aircraft is powered on for more than 51 days this can lead to **“display of misleading data”** to the pilots, with that data including airspeed, attitude, altitude and engine operating indications. On top of all that, the **stall warning horn and overspeed horn also stop working**. ...

The problem? They put **a millisecond clock with a 32-bit register and it overflows**.

- [https://twitter.com/mountain\\_ghosts/status/1245754158910705668](https://twitter.com/mountain_ghosts/status/1245754158910705668)
- [www.theregister.co.uk/2020/04/02/boeing\\_787\\_power\\_cycle\\_51\\_days\\_stale\\_data/](http://www.theregister.co.uk/2020/04/02/boeing_787_power_cycle_51_days_stale_data/)

## Real world example 4: Sequoia (2021)

*We discovered a **size\_t-to-int conversion vulnerability** in the Linux kernel's filesystem layer: ...*

*We **successfully exploited this uncontrolled out-of-bounds write, and obtained full root privileges** on default installations of Ubuntu 20.04, Ubuntu 20.10, Ubuntu 21.04, Debian 11, and Fedora 34 Workstation ...*

*To the best of our knowledge, this **vulnerability was introduced in July 2014** ...*

<https://www.qualys.com/2021/07/20/cve-2021-33909/sequoia-local-privilege-escalation-linux.txt>

## Real world example 5: FORCEDENTRY (2021)

... *a zero-day zero-click exploit against iMessage* ...

*All iPhones with iOS versions prior to 14.8, All Mac computers with operating system versions prior to OSX Big Sur 11.6, Security Update 2021-005 Catalina, and all Apple Watches prior to watchOS 7.6.2.*

...

*The exploit works by exploiting an integer overflow vulnerability in Apple's image rendering library (CoreGraphics). We are publishing limited technical information about CVE-2021-30860 at this time.*

<https://citizenlab.ca/2021/09/forcedentry-nso-group-imessage-zero-click-exploit-captured-in-the-wild/>



# A bad strcmp

A different example... what's wrong with this?

```
int bad_strcmp(const char *s1, const char *s2)
{
    while (*s1 && *s2 && *s1==*s2) {
        ++s1;
        ++s2;
    }
    return *s1 - *s2;
}
```

<https://pubs.opengroup.org/onlinepubs/009695399/functions/strcmp.html>

specifies that: "...the difference between the values of the first pair of bytes (both interpreted as type unsigned char) that differ..."; however, the C standard [ANS11] says that: "The three types char, signed char, and unsigned char [...] The implementation shall define char to have the same range, representation, and behavior as either signed char or unsigned char."

# SQL injection example (1/2)

```
# get the input from the user
name = ...
surname = ...

# build the SQL insert command
# eg.: INSERT INTO Students(name, surname) VALUES('foo', 'bar');
sql_command = "INSERT INTO Students(name, surname) VALUES('" +
               name + "', '" + surname + "');"

# ... and send the command to the DB server
execute_sql(sql_command)
```


# SQL injection example (2/2)

Consider:

```
name = "foo', 'bar'); DROP Students; -- "  
surname = "sql injection example"
```

Then, the command becomes:

```
"INSERT INTO Students(name, surname) VALUES('foo', 'bar');  
DROP Students; -- ', 'sql injection example');"
```



<https://xkcd.com/327/>

# The underlying problem

mixing code and (user provided) data; to avoid these problems you can:

- sanitize/escape data
  - use standard functions, don't roll your own, and
  - hope they are correct (e.g. `htmlspecialchars` vulnerability <https://www.cvedetails.com/cve/CVE-2009-4142/>)
- use compiled queries
- use domain-specific libraries/frameworks; e.g. ORMs
  - EF <https://msdn.microsoft.com/en-us/data/ef.aspx>
  - SQL Alchemy <https://www.sqlalchemy.org/>
  - ...

You can also use `sqlmap` <https://sqlmap.org/> on your application to check for vulnerabilities

# Command injection

```
import os

while True:
    f = input('Enter the filter (none to list all processes): ')
    cmd = 'ps aux'
    if f:
        cmd += " | grep '{}'.format(f)
    os.system(cmd)
```

# Log4Shell — CVE-2021-44228

- an **injection vulnerability** in Log4j, a popular logging framework
- the vulnerability had **existed unnoticed since 2013 (!)**
- simple exploit, **estimated to affect hundreds of millions of devices**
  - in the default configuration, when logging a string, Log4j performed string substitution on expressions of the form `${prefix:name}`
    - one recognized expression is `${jndi:lookup}`, where an arbitrary URL may be queried and loaded as Java object data; e.g.,  
`${jndi:ldap://example.com/file}`
  - by inputting a string that is logged, **an attacker could load and execute malicious code** hosted on a public URL
  - log4j 2.15.0 disabled this behavior by default
  - 2.16.0 completely removed this functionality

<https://en.wikipedia.org/wiki/Log4Shell>

# Stranger Strings — CVE-2022-35737

*SQLite is used in nearly everything, from naval warships to smartphones to other programming languages. The open-source database engine has a long history of being very secure. . .*

*SQLite implements custom versions of the printf family of functions and adds the new format specifiers %Q, %q, and %w, which are designed to properly escape quote characters in the input string in order to make safe SQL queries . . .*

*. . . is exploitable when large string inputs are passed to . . . SQLite . . . printf functions . . . cause the program to crash . . . if the format string contains the ! special character to enable unicode character scanning, then it is possible to achieve arbitrary code execution in the worst case, or to cause the program to hang . . .*

<https://blog.trailofbits.com/2022/10/25/sqlite-vulnerability-july-2022-library-api/>

- *all untrusted input* source must be validated/sanitized
  - ASAP!
- prefer allow-listing w.r.t deny-listing
  - deny-lists can be useful for testing; think of Apple's “goto fail”



# goto fail (2014 — Apple's SSL/TLS implementation)

```
int SSLVerifySignedServerKeyExchange(...)
{
    ...
    err = 0;
    ...
    if ((err = ReadyHash(&SSLHashSHA1, &hashCtx)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &clientRandom)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
        goto fail;
    goto fail;
    if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
        goto fail;
    ...
fail:
    ...
    return err;
}
```

<https://dwheeler.com/essays/apple-goto-fail.html>


# Memory errors

*Memory corruption bugs in software written in low-level languages like C or C++ are one of the oldest problems in computer security. The lack of safety in these languages allows attackers to alter the program's behavior or take full control over it by hijacking its control flow. This problem has existed for more than 30 years and a vast number of potential solutions have been proposed, yet memory corruption attacks continue to pose a serious threat. Real world exploits show that all currently deployed protections can be defeated. [SPWS13]*

The most infamous bug is. . . the buffer overflow

# What is a buffer overflow? (no technicalities today)

Exploiting a programming error, like this ☺:



October 22, 1895; [en.wikipedia.org/wiki/Montparnasse\\_derailment](https://en.wikipedia.org/wiki/Montparnasse_derailment) (idea:[Whe16])

# Read overflow

Rather than writing past the end of a buffer, a bug could permit *reading* past the end, **leaking information**

## Hearthbleed (2014)

Heartbleed is a security bug disclosed in April 2014 in the OpenSSL cryptography library, which is a widely used implementation of the Transport Layer Security (TLS) protocol.

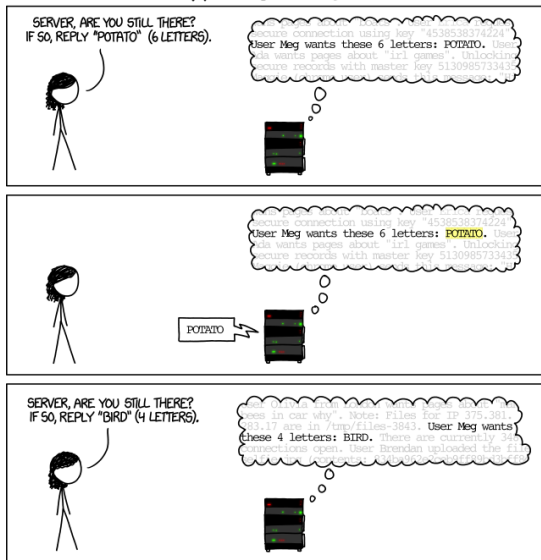
At the time of disclosure, some **17% (around half a million)** of the Internet's secure web servers certified by trusted authorities were believed to be vulnerable to the attack, allowing **theft of the servers' private keys and users' session cookies and passwords**.

<https://heartbleed.com/>

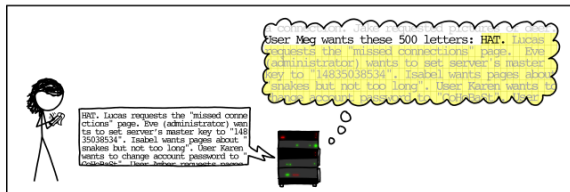

# Heartbleed (1/2)

<https://xkcd.com/1354/>

## HOW THE HEARTBLEED BUG WORKS:



<https://xkcd.com/1354/>



# What are the “inputs” of a program?

What can cause a change in a program behavior?

- Obvious ones
  - Command-line arguments
  - What the user types/clicks/...
  - File contents
- Not so obvious one: the **environment** (the program runs in)
  - the current directory
  - PATH
  - LD\_LIBRARY\_PATH
  - the inherited handles (file descriptors, sockets, ...)

E.g., <https://www.wietzebeukema.nl/blog/save-the-environment-variables>

*By manipulating environment variables on process level, it is possible to let trusted applications load arbitrary DLLs and execute malicious code. This post lists nearly 100 executables vulnerable to this type of DLL Hijacking on Windows 11 (21H2)*

See also: <https://hijacklibs.net/>

# Nebula 01

→ `examples/nebula1.c`

`https://exploit.education/nebula/level-01/`  
Level usernames *and* passwords: `levelxx`



# TOCTTOU (1/2)

Consider the following snippet:

- assume it is contained in a setuid (root) program
- note the use of `access`, to check whether the real user would be allowed to write the file

```
if (access("file", W_OK) != 0)
    exit(EXIT_FAILURE);
fd = open("file", O_WRONLY);
write(fd, buffer, sizeof(buffer));
```

- do you see any problem?
- does `access` really use the *real* user for checking?

## TOCTTOU (2/2)

**Time of check to time of use** is a class of bugs caused by changes in a system **between the checking** of a condition (such as a security credential) **and the use** of the results of such a check

```
/* victim */  
if (access("file", W_OK) != 0)  
    exit(EXIT_FAILURE);  
  
fd = open("file", O_WRONLY);  
write(fd, buffer, sizeof(buffer));
```

```
/* attacker */  
  
symlink("/etc/passwd",  
        "file");
```

Since Linux 3.6, symlink protection enabled by default

See `/proc/sys/fs/protected_symlinks` in `proc(5)`

# Dirty Cow

## Dirty COW

A race condition was found in the way the Linux kernel's memory subsystem handled the copy-on-write (COW) breakage of private read-only memory mappings. An unprivileged local user could use this flaw to gain write access to otherwise read-only memory mappings and thus increase their privileges on the system.



**DIRTY COW**

<https://dirtycow.ninja>

Existed in the Linux kernel since version 2.6.22 released 2007, and there is information about it being actively exploited at least since October 2016

# Example: Compilation service

Consider the following situation:

- A program provides **compilation services** to other programs
- The **client** program **specifies** the name of the **input and output files**
- the server is given the same access to those files that the client has
- however, the compiler service is **pay-per-use**, and the compiler service **stores its billing information** in a file (dubbed BILL) that **only it has access to**

Does it seem reasonable?

- Suppose a client calls the service and names its output file BILL
- The service opens the output file. Even though the client did not have access to that file, the service does, so the open succeeds, and the server writes the compilation output to the file, overwriting it, and thus destroying the billing information [Har88]

[https://en.wikipedia.org/wiki/Confused\\_deputy\\_problem](https://en.wikipedia.org/wiki/Confused_deputy_problem)

# Principle of the least privilege

## Principle of the least privilege

*Every program and every privileged user of the system should operate using the least amount of privilege necessary to complete the job.*

*Jerome Saltzer, Communications of the ACM*

The oldest instance of least privilege is probably the source code of `login.c`

- begins execution with super-user permissions and
- the instant they are no longer necessary, dismisses them via `setuid` with a non-zero argument
- as demonstrated in the Version 6 Unix source code:

<https://www.retro11.de/ouxr/u6ed/usr/source/s1/login.c.html#n:132>

From Wikipedia:

[https://en.wikipedia.org/wiki/Principle\\_of\\_least\\_privilege](https://en.wikipedia.org/wiki/Principle_of_least_privilege)

# Principle of the least privilege in OSes

Can you think of applications of this principle in modern OSes?

**Linux** **sudo** <https://en.wikipedia.org/wiki/Sudo>

Unfortunately, dropping unneeded privileges is error-prone, and there are portability issues [TDSW08]; moreover, we may hope sudo to be bug-free, but various bugs have been found (and fixed).

See also, *Zero-day vulnerability in Bash - Suidbash Google CTF Finals 2019*

<https://www.youtube.com/watch?v=-wGtxJ8opa8>

**Windows** **UAC (User Account Control)**

[https://en.wikipedia.org/wiki/User\\_Account\\_Control](https://en.wikipedia.org/wiki/User_Account_Control)

unfortunately, **useless with default settings (!)** see

<https://github.com/hfiref0x/UACME> and

[devblogs.microsoft.com/oldnewthing/20160816-00/?p=94105](https://devblogs.microsoft.com/oldnewthing/20160816-00/?p=94105)

# Format string vulnerabilities

Just the idea; see [stt01, HLV10] for more

```
#include <stdio.h>

int main(int argc, char* argv[])
{
    int i;
    for(i = 1; i < argc; ++i)
        printf(argv[i]);
    printf("\n");
}
```

Think of %x, %p, or %s; moreover... and we'll discuss also %n in the next lectures

# Bad strcmp revisited

```
int still_bad_strcmp(const char *s1, const char *s2)
{
    while (*s1 && *s2 && *s1==*s2) {
        ++s1;
        ++s2;
    }
    return *((unsigned char *)s1) - *((unsigned char *)s2);
    // (unsigned) chars are promoted
    // to int to perform the subtraction
}

bool check_password(const char *password)
{
    return still_bad_strcmp(password, "zxgio") == 0;
}
```

...what's wrong here?



# Side channels (1/2)

## Side channel attacks

In computer security, a side-channel attack is any attack based on **information gained from the implementation** of a computer system, rather than weaknesses in the implemented algorithm itself (e.g. cryptanalysis and software bugs). **Timing information, power consumption, electromagnetic leaks or even sound** can provide an extra source of information, which can be exploited.

[https://en.wikipedia.org/wiki/Side-channel\\_attack](https://en.wikipedia.org/wiki/Side-channel_attack)

## Side channels (2/2)

Also:

*Side channel = “Obtaining meta-data and deriving data from it.”  
by Daniel Gruss (@lavados)*

Thread:

<https://twitter.com/lavados/status/1156982866414379008>

Important recent examples are Meltdown [LSG<sup>+</sup>18], Spectre [KHF<sup>+</sup>19] and Bleichenbacher’s CAT [RGG<sup>+</sup>18]

# Outline

## 1 Introduction

## 2 Common Weaknesses

- Integer overflows (& undefined/implementation-defined behaviours)
- Command injection
- Memory errors
- “Hidden inputs”
- Race conditions
- Privileges
- Information leakage

## 3 Secure Coding

- Static Analysis
- Dynamic Analysis
  - Sanitizers
  - Fuzzers

How can we **find/avoid vulnerabilities**?

There are many techniques and tools, with no clear winners between:

- **Static Analyses**

- Can work on incomplete code
- A single run can analyze all code
- However, there are false positives and false negatives

- **Dynamic Analyses**

- Usually, easy to set up and run
- No false positives
- No false negatives *but* limited to what has been executed
- Slows down execution

# Static analysis tools

- behave a bit like spell-checkers
  - they prevent well-understood varieties of mistakes from going unnoticed
  - a clean run doesn't guarantee that code is perfect
    - it is probably just free of certain kinds of common problems
  - however, they don't automatically make you an excellent coder 😊
- can find errors early in development, even before the first run
- can recheck large bodies of code when a new attack is discovered

# No silver bullets

The problem with *interesting* static analyses is that they are **undecidable**; in practice, most “work” (=produce useful results) but:

- **false positives**/alarms
- **false negatives** — unreported problems that exist in the program

The balance between them is often indicative of the purpose of the tool:

- code quality tools usually produce a low number of false positives
- security tools usually produce more false positives

## Static analysis tools check the code (only)

To catch a defect, it must be “visible” in the code: architectural risk analysis is a necessary complement

# Type checking

In Java, for instance, we have both **false positives**:

```
short s = 0;
int i = s;
short r = i; // error: incompatible types: possible lossy
              // conversion from int to short
```

And **false negatives**:

```
Object[] objs = new String[1];
objs[0] = new Object(); // Exception in thread "main"
                        // java.lang.ArrayStoreException:
                        // java.lang.Object
```

# Program verification

Program verification tools try to prove that some code is a **faithful implementation of a specification**

- creating proper specifications can require more work than writing code
- historically, these tools could not process programs of significant size

More commonly, verification tools **check software against a *partial specification*** that details only part of the behavior of a program. This is sometimes called **property checking**

For instance, ...



# Property checking example

For instance, properly allocating/releasing memory:

```
int f()
{
    char *inBuf, *outBuf;
    inBuf = malloc(512);
    outBuf = malloc(512);
    if (outBuf == 0 || inBuf == 0)
        return -1; // [...] A memory leak is possible.
```

Ok; and now?

```
int f()
{
    char *inBuf, *outBuf;
    inBuf = malloc(512);
    if (inBuf == 0)
        return -1; // as before... (for some tools)
```

A middle ground between style-checking and program verification

A bug finder points out places where the program will behave in a way that the programmer did not *probably* intend


- most come pre-stocked with a set of “bug idioms” (rules) that describe patterns that often indicate bugs

E.g. CLang Static Analyzer or FindBugs (Java)

# Inner working

How do these tool work? The first thing is transforming the code to be analyzed into a **model**, that is, a set of data structures

- the kind of model depends on the analysis
- many data-structures and algorithms shared with **compiler world**



From [CW07]

# Reporting results

If you can't make sense of what a tool reports, the result is useless; e.g.

- too many false positives
- bad presentation of “good” results can be confused with analysis mistakes

tools should offer

- integration with IDEs
- a way to *easily*
  - navigate, group and sort results (by severity, confidence, ...)
  - eliminate/suppress unwanted warnings
  - explain the reason for each warning
  - provide recommendations

# Crucial issue: context sensitivity

**Context sensitivity**, that is, circumstances and conditions under which a particular piece of code runs is crucial:

- Easy to point at all calls to `strcpy`
- Hard to pinpoint specific calls that might allow an attacker to overflow a buffer

E.g.

```
int main(int argc, char **argv)
{
    char buf1[1024];
    char buf2[1024];
    strcpy(buf1, "pizza");
    strcpy(buf2, argv[0]);
    ...
}
```

# How to use static-analysis tools

Some tools are simply run like compilers

- i.e , you run them on source files and get the list of warnings in return

Others, require a multi-step approach

- a command “monitors” the building process, producing either
  - the analysis result, or
  - some intermediate files; that are then processed by another command, the *analyzer*
- another command shows the results or export them in a (PDF/HTML/...) report

# CppCheck

Cppcheck is a static analysis tool for C/C++ code. It focuses on detecting undefined behaviour and dangerous coding constructs. The goal is to detect only real errors in the code (i.e. have very few false positives).

<https://cppcheck.sourceforge.net/>

- Installation: `sudo apt install cppcheck`
- Usage: `cppcheck file-or-path`

```
kill.c:117:11: style: Local variable 'rv' shadows outer variable [shadowVariable]
    int rv = print_signal_list();
    ^
kill.c:106:7: note: Shadowed declaration
    int rv = 0;
    ^
kill.c:117:11: note: Shadow variable
    int rv = print_signal_list();
    ^
shell.c:106:27: style: The scope of the variable 'prev' can be reduced. [variableScope]
    procedure *cur = proc, *prev;
                        ^
tokenizer.c:96:17: style: The scope of the variable 'repl' can be reduced. [variableScope]
    char* repl;
        ^
```

Clang Static Analyzer is a source code analysis tool that finds bugs in C, C++, and Objective-C programs

`https://clang-analyzer.llvm.org`

- Installation: `sudo apt install clang-tools`
- Usage:
  - `scan-build` monitors the build process
  - `scan-view` creates a web-server that allows you to browse the issues

For instance, ...



# Clang Static Analyzer example (1/2)

## Bug Summary

Bug Type	Quantity	Display?
<b>All Bugs</b>	<b>17</b>	<input checked="" type="checkbox"/>
<b>Dead store</b>		
Dead assignment	1	<input checked="" type="checkbox"/>
<b>Logic error</b>		
Dereference of undefined pointer value	2	<input checked="" type="checkbox"/>
<b>Memory error</b>		
Double free	2	<input checked="" type="checkbox"/>
Memory leak	5	<input checked="" type="checkbox"/>
<b>Security</b>		
Potential insecure memory buffer bounds restriction in call 'strcat'	2	<input checked="" type="checkbox"/>
Potential insecure memory buffer bounds restriction in call 'strcpy'	5	<input checked="" type="checkbox"/>

## Reports

Bug Group	Bug Type ▼	File	Function/Method	Line	Path Length			
Dead store	Dead assignment	tokenizer.c	tokenize	195	1	<a href="#">View Report</a>	<a href="#">Report Bug</a>	<a href="#">Open</a>
Logic error	Dereference of undefined pointer value	shell.c	build_procedure_list	90	20	<a href="#">View Report</a>	<a href="#">Report Bug</a>	<a href="#">Open</a>
Logic error	Dereference of undefined pointer value	shell.c	build_procedure_list	93	17	<a href="#">View Report</a>	<a href="#">Report Bug</a>	<a href="#">Open</a>
Memory error	Double free	tokenizer.c	tokenize	182	38	<a href="#">View Report</a>	<a href="#">Report Bug</a>	<a href="#">Open</a>
Memory error	Double free	tokenizer.c	tokenize	115	25	<a href="#">View Report</a>	<a href="#">Report Bug</a>	<a href="#">Open</a>

# Clang Static Analyzer example (2/2)

```
169
170
171
172
173
else{
    char* envVar = strdup(&line[i+1], j-i-1);
    char* variable_value = getenv(envVar);

    if(variable_value)
```

31 ← Assuming 'variable\_value' is non-null →

32 ← Taking true branch →

```
174
175
176
177
178
179
180
    repl = variable_value;
    else
        repl = getsudoenv(envi, envVar);

    strcpy(&token[n], repl);
    n += strlen(repl);
    free(repl);
```

33 ← Memory is released →

```
181
182
    free(envVar);
    free(variable_value);
```

34 ← Attempt to free released memory

PVS-Studio is a tool for detecting bugs and security weaknesses in the source code of programs, written in C, C++, C# and Java.

- Commercial, but free for students:  
<https://www.viva64.com/en/for-students/>
- Usage:
  - `pvs-studio-analyzer trace` monitors the build process
  - `pvs-studio-analyzer analyze` analyzes 😊
  - `plog-converter` exports in various formats

For instance, ...

# PVS-Studio example (1/2)



## PVS-Studio Analysis Results

<b>Date:</b>	Thu Nov 19 12:25:33 2020
<b>PVS-Studio Version:</b>	7.10.43305.85
<b>Command Line:</b>	plog-converter -t fullhtml -a GA\64\OP\CS shell.log -o pvs-shell-analysis
<b>Total Warnings (GA):</b>	27

Group	Location	Level	Code	
General Analysis	<a href="#">execute.c:62</a>	Medium	<a href="#">V522</a>	Dereferencing of the null pointer 'path' might take place. The potential null
General Analysis	<a href="#">execute.c:363</a>	Medium	<a href="#">V575</a>	The potential null pointer is passed into 'strchr' function. Inspect the first a
General Analysis	<a href="#">export.c:11</a>	Medium	<a href="#">V575</a>	The potential null pointer is passed into 'strchr' function. Inspect the first a
General Analysis	<a href="#">pwd.c:9</a>	Medium	<a href="#">V575</a>	The potential null pointer is passed into 'getcwd' function. Inspect the first
General Analysis	<a href="#">pwd.c:12</a>	High	<a href="#">V575</a>	The null pointer is passed into 'free' function. Inspect the first argument.
General Analysis	<a href="#">shell.c:87</a>	Medium	<a href="#">V522</a>	There might be dereferencing of a potential null pointer 'cur'. Check lines:
General Analysis	<a href="#">shell.c:128</a>	Medium	<a href="#">V755</a>	A copy from unsafe data source to a buffer of fixed size. Buffer overflow is
General Analysis	<a href="#">shell.c:136</a>	Medium	<a href="#">V755</a>	A copy from unsafe data source to a buffer of fixed size. Buffer overflow is
General Analysis	<a href="#">shell.c:148</a>	Medium	<a href="#">V728</a>	An excessive check can be simplified. The '  ' operator is surrounded by o
General Analysis	<a href="#">sudo_environment.c:19</a>	Medium	<a href="#">V701</a>	realloc() possible leak: when realloc() fails in allocating memory, original p
General Analysis	<a href="#">sudo_environment.c:23</a>	Low	<a href="#">V522</a>	There might be dereferencing of a potential null pointer 'e->elems'.

# PVS-Studio example (2/2)

```
117 int main(unused int argc, unused char *argv[]) {
118     init_shell();
119
120     static char line[4096];
121     int line_num = 0;
122
123     int loop = true;
124
125     if(argc > 1){
126         if(strcmp(argv[1], "-c") == 0){
127             loop = false;
128             strcpy(line, argv[2]);
129         }else{
130             fprintf(stderr, "wrong flag");
131         }
132     }
```

↑ [V755](#) A copy from unsafe data source to a buffer of fixed size. Buffer overflow is possible.

# Facebook Infer

*Infer is a static analysis tool - if you give Infer some **Java** or **C/C++/Objective-C** code it produces a list of potential bugs.*

<https://fbinfer.com/>  
<https://github.com/facebook/infer/releases>

For instance, with this program:

```
#include <stdlib.h>

void test() {
    int *s = NULL;
    *s = 42;
}
```

you get:

```
hello.c:5: error: NULL_DEREFERENCE
    pointer s last assigned on line 4 could be null and is dereferenced at line 5
```

# Dynamic code analysis

Dynamic analysis tools work by

- 1 (statically or dynamically) **instrumenting** the target application
- 2 **observing** an execution to detect errors

Major players: Valgrind, Dr.Memory and clang/gcc sanitizers

# Valgrind

Instrumentation framework for building dynamic analysis tools, can detect many **memory** management and **threading** bugs, and profile your programs

- Works on **Unix-like OSes** (i.e. no Windows)
- Runs **unmodified binaries**; however, it's better to
  - compile with symbols (`-g`) and
  - without/with-little optimizations; otherwise, valgrind occasionally reports spurious errors
- Only for heap allocations (i.e. no stack/global)
- **Huge slow-down** (20-30x)
- Installation: `apt install valgrind`
- Quick-start: compile with debug symbols, then
  - `valgrind [--leak-check=full] executable args` # memory
  - `valgrind --tool=helgrind executable args` # threading
- Offers a gdb-server:
  - `valgrind --vgdb=full --vgdb-error=0 executable args`

<https://valgrind.org>



Similarly to Valgrind, **identifies memory-related programming errors**

- Works on Windows, Linux (+Android), Mac
- Runs unmodified binaries (but same considerations as valgrind)
- Generally faster than Valgrind
- Distributed as a tar.gz
- Quick-start:

```
g++ -m32 -g -fno-inline -fno-omit-frame-pointer ...  
drmemory -- executable args
```

<https://drmemory.org>

# Outline

## 1 Introduction

## 2 Common Weaknesses

- Integer overflows (& undefined/implementation-defined behaviours)
- Command injection
- Memory errors
- “Hidden inputs”
- Race conditions
- Privileges
- Information leakage

## 3 Secure Coding

- Static Analysis
- **Dynamic Analysis**
  - **Sanitizers**
  - Fuzzers

# Sanitizers

We'll consider LLVM Sanitizers [SBPV12]; for a full picture see also [SLR<sup>+</sup>18]

- use compile-time instrumentation (Clang, some work in gcc too)
- detect **memory**, **thread** errors and **undefined behaviors**
- way faster than Valgrind/Dr.Memory

they are:

- AddressSanitizer (+LeakSanitizer) — addressability issues and memory leaks
  - On some platform there is also a more efficient Hardware-assisted AddressSanitizer
- MemorySanitizer — use of uninitialized memory
- ThreadSanitizer — data races and deadlocks
- UndefinedBehaviorSanitizer – (some) undefined behaviors

# Address Sanitizer

## Finds

- buffer overflows (stack, heap, globals)
- heap-use-after-free, stack-use-after-return
- leaks, init-order, double-free, ...

## by

- instrumenting all loads/stores, and replacing malloc and friends

## To use:

```
-fsanitize=address -g -fno-common -fno-omit-frame-pointer
```

## Beware:

- ① `-fsanitize=address` must be used both when compiling and linking; i.e. `C(PP)FLAGS` and `LDFLAGS`
- ② uses ptrace: use `ASAN_OPTIONS=detect_leaks=0` `gdb ...` to debug (don't export, otherwise detection would be always disabled)

Special checks can be enabled via `ASAN_OPTIONS`; see

<https://github.com/google/sanitizers/wiki/AddressSanitizerFlags>

# How does this work?

The instrumentation replaces (de)allocation functions, that

- **insert redzones** around every allocation, as we'll discuss
- **delay the reuse** of freed memory
  - poisoning the entire memory region on free
- **collect stack traces** for every malloc/free

## Overhead

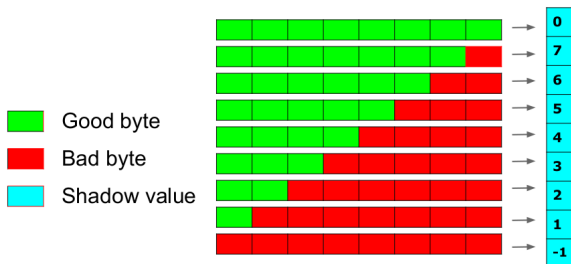
- 2x slowdown
- 1.5x-3x memory

# Shadow map

Part of the address space is reserved for a **shadow map**, that encodes the state of the other parts

Any aligned 8 bytes may have 9 states:

N good bytes and 8 - N bad ( $0 \leq N \leq 8$ )



# Instrumentation

```
uint64_t *a = ...  
*a = ... // 8-byte access
```



```
char *shadow = (a >> 3) + shadowOffset;  
if (*shadow)  
    ReportError(a);  
*a = ...
```

accessing 1, 2, 4 bytes is a bit more involved (same idea, though)

# Instrumenting stack frames

```
void foo() {  
    char a[328];  
    <----- CODE ----->  
}
```




```
void foo() {  
    char rz1[32]; // 32-byte aligned  
    char a[328];  
    char rz2[24];  
    char rz3[32];  
    int *shadow = (&rz1 >> 3) + shadowOffset;  
    shadow[ 0] = 0xffffffff; // poison rz1  
    shadow[11] = 0xffffffff00; // poison rz2  
    shadow[12] = 0xffffffff; // poison rz3  
    <----- CODE ----->  
    shadow[0] = shadow[11] = shadow[12] = 0;  
}
```



# HW-assisted Address Sanitizer (1/2)

*HWASan is based on memory tagging ... Every memory allocation is assigned a random 8-bit tag that is stored in the most significant byte (MSB) of the address, but ignored by the CPU*



Better granularity and performance (especially when we'll have fully hardware-supported memory tagging, as in ARM v8.5)

# HW-assisted Address Sanitizer (2/2)

## Resources:

- **Memory Tagging for the Kernel: Tag-Based KASAN** by Andrey Konovalov @ Android Security Symposium 2020  
<https://www.youtube.com/watch?v=f-Rm7JFsJGI>
- **Hardware-Assisted Address Sanitizer (HWASan) in Android**  
<https://android-developers.googleblog.com/2020/02/detecting-memory-corruption-bugs-with-hwasan.html>
- **Memory Tagging and how it improves C/C++ memory safety** by Kostya Serebryany @ CppCon 2018  
<https://www.youtube.com/watch?v=1LEcbXidK2o>

# Memory Sanitizer

Detects uninitialized memory reads

Idea:

- Bit to bit shadow mapping (1 means *poisoned*; that is, uninitialized)
  - e.g., if `foo` is uninitialized, then `{foo &= 1;}` zero-initializes all its bits except for the least significant one!
- Memory returned by `malloc` and stack objects are *uninitialized*
- Shadow is unpoisoned when constants are stored
- MSan requires to recompile all libraries (to avoid false positives)
  - Libc can be wrapped, but inline-assembly and JIT?

`-fsanitize=memory`

`-fsanitize=memory -fsanitize=memory-track-origins`

## Overhead

- Without origins: 2.5x slowdown, 2x memory
- With origins: 5x slowdown, 3x memory

# Thread Sanitizer

Detects data races and deadlocks

- Compile-time instrumentation (LLVM, GCC)
  - Intercepts all reads/writes
- Run-time library
  - Replaces/intercept memory/synchronization functions

`-fsanitize=thread`

## Overhead/limitations

- Only 64-bit Linux, does not instrument libraries and inline assembly
- 4x-10x slowdown (still way faster than helgrind), 5x-8x memory overhead

# Undefined-behaviour Sanitizer

Modifies the program at compile-time, to catch various kinds of undefined behavior:

- Using misaligned or null pointer
- Signed integer overflow
- Conversion to, from, or between floating-point types which would overflow the destination
- ...

See the list of available checks at

<https://clang.llvm.org/docs/UndefinedBehaviorSanitizer.html>

`-fsanitize=undefined`

## Overhead

- 0 – 0.5x slowdown

# Outline

## 1 Introduction

## 2 Common Weaknesses

- Integer overflows (& undefined/implementation-defined behaviours)
- Command injection
- Memory errors
- “Hidden inputs”
- Race conditions
- Privileges
- Information leakage

## 3 Secure Coding

- Static Analysis
- Dynamic Analysis
  - Sanitizers
  - Fuzzers

# Introduction

*It was a dark and stormy night. Really. Sitting in my apartment in Madison in the Fall of 1988, there was a wild midwest thunderstorm pouring rain and lighting up the late night sky. That night, I was **logged on to the Unix systems** in my office **via a dial-up phone line** over a 1200 baud modem. With the heavy rain, there was noise on the line and that noise was interfering with my ability to type sensible commands to the shell and programs that I was running. It was a race to type an input line before the noise overwhelmed the command. This fighting with the noisy phone line was not surprising. What did surprise me was the fact that **the noise seemed to be causing programs to crash**. And more surprising to me was the programs that were crashing—common Unix utilities that we all use everyday.*

Barton Miller

from the book “Fuzzing for Software Security Testing and Quality Assurance”

# History

The term **fuzzing** originates from those days, and those ideas generated a stream of research, e.g. [MFS90]; however:

- Testing programs with random inputs dates back to the 1950s, when data was still stored on punched cards
- The execution of random inputs is also called **random testing** or **monkey testing**

## Generating Software Tests – Breaking Software for Fun and Profit

Very interesting resource, a “Textbook for Paper, Screen, and Keyboard” can be found at: <https://www.fuzzingbook.org/>



# Fuzz-testing, AKA Fuzzing

- Basic idea: throwing “random garbage” into programs and make them **crash** ... in interesting and **different ways**
  - sanitizers can be very useful
- How to generate “garbage”? More importantly, “quality garbage”?
  - Throw a coin, repeatedly 😊
    - Randomness is ok, totally random inputs may be useless
    - Inputs should be “representative” of expected format
  - Generate random test-case from a model; e.g. from grammars [GZ19]
  - Randomly mutate legal inputs; e.g. Radamsa  
<https://gitlab.com/akihe/radamsa>
  - **Leverage code-coverage**: e.g. using genetic algorithms to automatically discover clean, interesting test **cases that trigger new internal states** in the binary → **AFL**

# American Fuzzy Lop (fuzzer)

AFL by Michal Zalewski

- is a security-oriented **fuzzer**
- employs **compile-time instrumentation**
  - on Linux, optional QEMU mode allows black-box binaries to be fuzzed
- uses **genetic algorithms to discover interesting test cases**, that trigger new internal states in the targeted binary


# (Simplified) overall algorithm

- 1 load user-supplied initial test cases into the queue
- 2 take next input file from the queue
- 3 attempt to trim the test case to the smallest size that doesn't alter the measured behavior of the program
- 4 repeatedly mutate the file using a balanced and well-researched variety of traditional fuzzing strategies
- 5 if any of the generated mutations resulted in a new state transition recorded by the instrumentation, add mutated output as a new entry in the queue
- 6 go to (2)

The discovered test cases are also periodically culled to eliminate ones that have been obsoleted by newer, higher-coverage finds

# Pulling JPEGs out of thin air

*... created a text file containing just "hello" and asked the fuzzer to keep feeding it to a program that expects a JPEG image ... The first image, hit after about six hours on an 8-core system, looks very unassuming ... But the moment it is discovered, the fuzzer starts using the image as a seed - rapidly producing a wide array of more interesting pics ...*



<https://lcamtuf.blogspot.it/2014/11/pulling-jpegs-out-of-thin-air.html>

# How to get AFL

- Official repository: <https://github.com/Google/afl>
- **AFL++ [FMEH20] is a superior fork** to Google's AFL: more speed, more and better mutations, more and better instrumentation, custom module support, etc.

<https://github.com/AFLplusplus/AFLplusplus>



# References I

- [ANS11] ANSI/ISO.  
Working draft, Standard for Programming Language C.  
Technical Report N1570, ANSI/ISO, 2011.
- [ANS12] ANSI/ISO.  
Working draft, Standard for Programming Language C++.  
Technical Report N3337, ANSI/ISO, 2012.
- [CW07] Brian Chess and Jacob West.  
*Secure programming with static analysis*.  
Pearson Education, 2007.
- [DMS06] Mark Dowd, John McDonald, and Justin Schuh.  
*The art of software security assessment: Identifying and preventing software vulnerabilities*.  
Pearson Education, 2006.

- [FMEH20] Andrea Fioraldi, Dominik Maier, Heiko Eifeldt, and Marc Heuse.  
AFL++: Combining incremental steps of fuzzing research.  
*In 14th USENIX Workshop on Offensive Technologies (WOOT 20)*, 2020.
- [GZ19] Rahul Gopinath and Andreas Zeller.  
Building fast fuzzers, 2019.
- [Har88] Norm Hardy.  
The confused deputy (or why capabilities might have been invented).  
*ACM SIGOPS Operating Systems Review*, 22(4):36–38, 1988.

# References III

- [HLV10] Michael Howard, David LeBlanc, and John Viega.  
*24 Deadly Sins of Software Security: Programming Flaws and How to Fix Them.*  
McGraw-Hill, Inc., 2010.
- [KHF<sup>+</sup>19] Paul Kocher, Jann Horn, Anders Fogh, , Daniel Genkin, Daniel Gruss, Werner Haas, Mike Hamburg, Moritz Lipp, Stefan Mangard, Thomas Prescher, Michael Schwarz, and Yuval Yarom.  
Spectre attacks: Exploiting speculative execution.  
In *40th IEEE Symposium on Security and Privacy (S&P'19)*, 2019.
- [KP81] Brian W. Kernighan and P. J. Plauger.  
*Software Tools in Pascal.*  
Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 1st edition, 1981.



# References IV

- [LSG<sup>+</sup>18] Moritz Lipp, Michael Schwarz, Daniel Gruss, Thomas Prescher, Werner Haas, Anders Fogh, Jann Horn, Stefan Mangard, Paul Kocher, Daniel Genkin, Yuval Yarom, and Mike Hamburg.  
Meltdown: Reading kernel memory from user space.  
*In 27th USENIX Security Symposium (USENIX Security 18)*, 2018.
- [MFS90] Barton P Miller, Louis Fredriksen, and Bryan So.  
An empirical study of the reliability of unix utilities.  
*Communications of the ACM*, 33(12):32–44, 1990.
- [Pol17] Erik Poll.  
Lecture notes on language-based security, 2017.

- [RGG<sup>+</sup>18] Eyal Ronen, Robert Gillham, Daniel Genkin, Adi Shamir, David Wong, and Yuval Yarom.  
The 9 Lives of Bleichenbacher's CAT: New Cache ATtacks on TLS Implementations, 2018.  
<http://cat.eyalro.net/>.
- [SBPV12] Konstantin Serebryany, Derek Bruening, Alexander Potapenko, and Dmitriy Vyukov.  
AddressSanitizer: A Fast Address Sanity Checker.  
In *USENIX Annual Technical Conference*, pages 309–318, 2012.
- [SLR<sup>+</sup>18] Dokyung Song, Julian Lettner, Prabhu Rajasekaran, Yeoul Na, Stijn Volckaert, Per Larsen, and Michael Franz.  
SoK: Sanitizing for Security.  
*arXiv preprint arXiv:1806.04355*, 2018.

- [SPWS13] Laszlo Szekeres, Mathias Payer, Tao Wei, and Dawn Song.  
Sok: Eternal war in memory.  
In *Security and Privacy (SP), 2013 IEEE Symposium on*, pages 48–62. IEEE, 2013.
- [stt01] scut / team teso.  
Exploiting format string vulnerabilities.  
<https://crypto.stanford.edu/cs155/papers/formatstring-1.2.pdf>, 2001.
- [TDSW08] Dan Tsafir, Dilma Da Silva, and David Wagner.  
The murky issue of changing process identity: revising “setuid demystified”.  
*USENIX Login*, 33(3):55–66, 2008.

- [WCC<sup>+</sup>12] Xi Wang, Haogang Chen, Alvin Cheung, Zhihao Jia, Nickolai Zeldovich, and M Frans Kaashoek.  
Undefined behavior: what happened to my code?  
In *Proceedings of the Asia-Pacific Workshop on Systems*,  
page 9. ACM, 2012.
- [Whe16] David A. Wheeler.  
Secure programming HOWTO, 2016.  
<http://www.dwheeler.com/secure-programs/>.